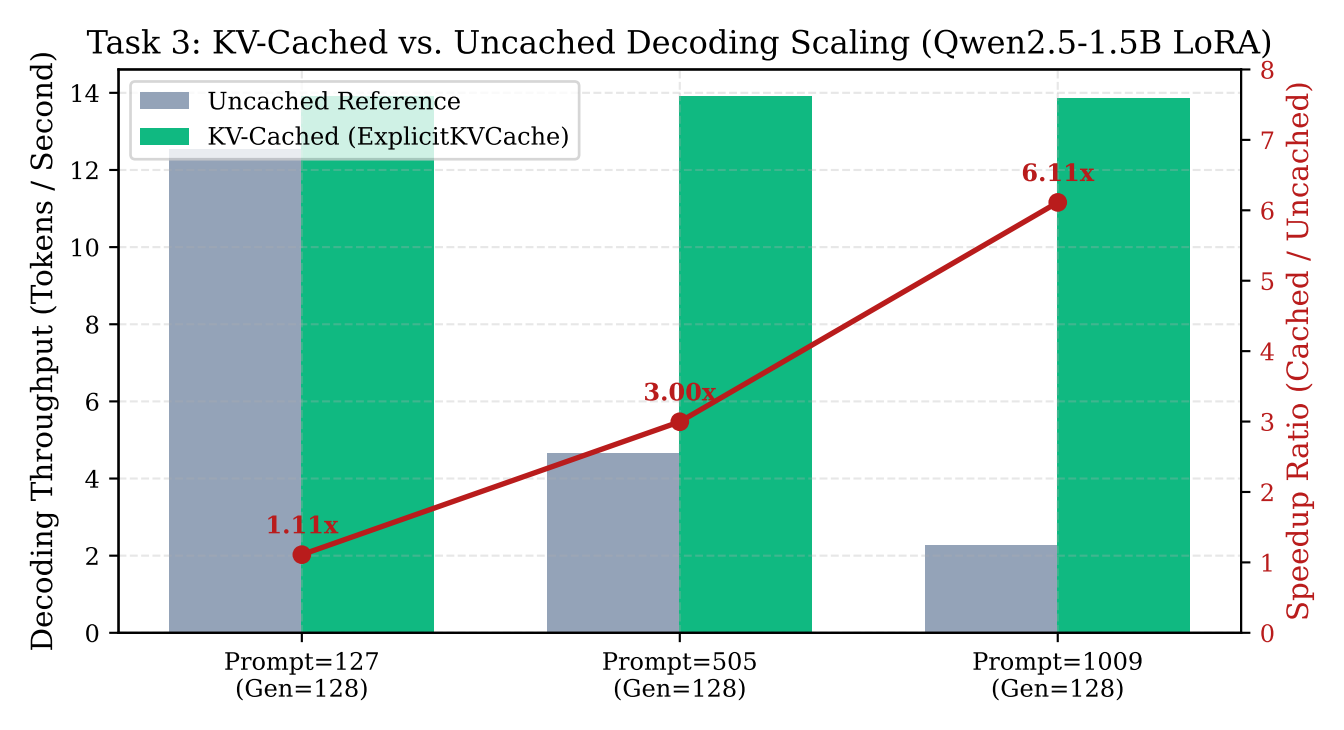
Task 3: KV-Cached vs. Uncached Decoding Scaling (Qwen2.5-1.5B LoRA)
Uncached Reference
KV-Cached (ExplicitKVCache)
Decoding Throughput (Tokens / Second)
Speedup Ratio (Cached / Uncached)
1.11x
3.00x
6.11x
Prompt=127
(Gen=128)
Prompt=505
(Gen=128)
Prompt=1009
(Gen=128)